



Inference Efficiency

Latency



Token Usage



Retry Overhead

